

MUSINGS OF AN IDIOT

Adrian Jackson

a.jackson@epcc.ed.ac.uk

@adrianjhpc





BE NICE

BE KIND

FIND A TEAM

FIND A BALANCE

- Be wrong
- Ask questions
- Work on things you enjoy
- Work in public
- Seniority or money isn't everything



PARALLEL COMPUTING

Adrian Jackson


adrianj@epcc.ed.ac.uk

@adrianjhpc



HPC - Parallelism

- Simulation science drives computing power
 - Consistently more computation power needed than available
 - Runtime of months on a single processor not uncommon
 - Parallel programs often start out as serial programs
- Why not just make a faster chip?
 - Theoretical problems
 - Physical limitations to size and speed of a single chip
 - Speed of light, size of atoms, dissipation of heat
 - Voltage reduction versus clock speed for power requirements
 - Voltages becomes too small for “digital” differences
 - Practical problems
 - Developing new chips is incredibly expensive
 - Must make maximum use of existing technology
- Solution
 - Use many CPUs cooperatively on the same problem
 - HPC computing synonymous with parallelism




Building Blocks of Parallel Machines

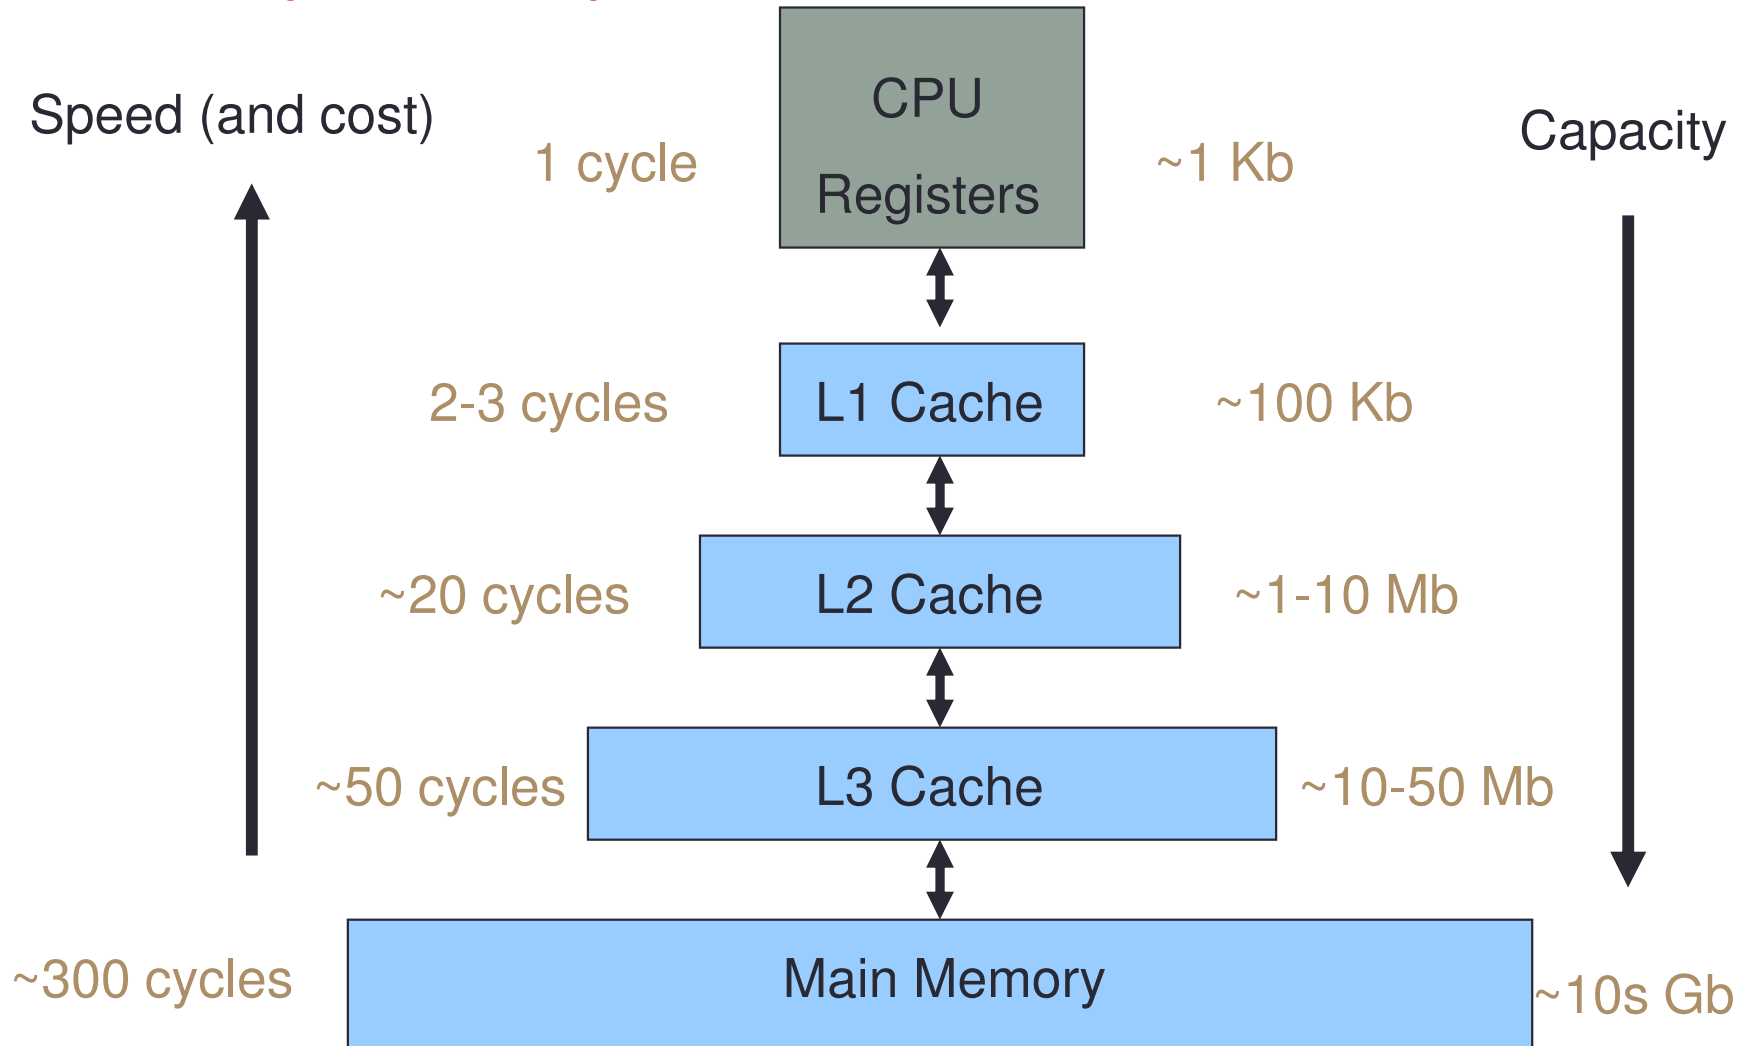
Four principal technologies which make up computers

- Processors
 - to calculate
- Memory
 - for temporary storage of data
- Interconnect
 - so processors can talk to each other (and the outside world)
- Storage
 - disks for storing input/output data and tapes for long term archiving of data

Performance Trend





Memory hierarchy



Types of HPC systems

- Shared-memory: OpenMP
 - Multiple processors share a single memory space
 - Simple to program for many problems
 - Scaling is problematic
- Distributed memory: MPI
 - Each processing unit has its own memory space
 - Excellent scaling properties
 - Can be more complex to program due to explicit communications
- Accelerators (GPU, Intel MIC)
 - Specialist processing units attached to main CPU
 - Can be difficult to extract good performance
 - (Conceptually similar to old vector architectures.)



Distributed Shared Memory (clusters)



- Dominant architecture is a hybrid of these two approaches: *Distributed Shared Memory*.
 - Due to most HPC systems being built from commodity hardware – trend to multicore processors.
 - Each Shared memory block is known as a *node*.
 - Usually 16-64 processors per node.
 - Nodes can also contain accelerators.
- Majority of users try to exploit in the same way as for a purely distributed machine
 - As the number of cores per node increases this can become increasingly inefficient...
 - ...and programming for these machines can become increasingly complex

Differences from Cloud computing

- Performance
 - Clouds usually use virtual machines which add an extra layer of software.
 - In cloud you often share hardware resource with other users – HPC access is usually exclusive.
- Tight-coupling
 - HPC parallel programming usually assumes that the separate processes are tightly coupled
 - Requires a low-latency, high-bandwidth communication system between tasks
 - Cloud usually does not usually have this
- Programming models
 - HPC use high-level compiled languages with extensive optimisation.
 - Cloud often based on interpreted/JIT.


New hardware trends

- Many-core
 - Intel KNL processors
 - 72 cores, 288 threads
 - Nvidia Pascal
 - Large scale GPU
- Memory and I/O becoming more complex
 - Different levels of memory
 - Stacked, DRAM, NVRAM
 - Different levels of I/O
 - NVRAM, SSDs, Lustre, /tmp
- More heterogeneity in both processing and memory systems
 - Applications are likely to have to adapt to get best performance




Usage by codes


Usage by programming language




Usage by code





What is it used for?



Institutional use



Machine Usage



What is a software developer?


- Coder
 - Writes code that compiles and runs
 - Inefficient, full of bugs, difficult to understand, fix and extend
- Programmer
 - Uses tools to develop correct code more easily, in less time, with less pain
 - Efficient, correct, modular, easy to understand, fix and extend
- Software developer
 - Gathers requirements
 - Creates designs and assesses alternatives
 - Writes, fixes, improves and extends programs
 - Writes user doc
 - Prepares releases
 - Trains users and provides support
- But what if my goal is to become a good researcher?

Programming skills and research


- Good programming skills are good research skills
- Record which version of your software produced the data you used in a paper, presentation, poster or thesis
 - Version control
- Help others to validate what you did by peer review, and to replicate, reproduce and reuse what you did
 - Good programming practice
- Ensure that your algorithms, implementation and data are correct and so that your research is correct
 - Unit testing, debugging, profiling
- Free up time for research
 - Automation and build tools, test frameworks

- Programming Language is irrelevant
- Learn at least 2 languages


Fully science workflow is important



Fully science workflow is important



Improved workflow



- 256 block grid, high load imbalance
 - 90% difference in grid points between largest and smallest blocks

Summary

- Programs run on computers (doh!)
- Understanding computers helps understanding programs
 - Running and writing
- Programming takes effort
 - Learn more than one programming language
 - The more perspectives you have the better you'll be

What now?

- You can attempt the ARCHER driving test
 - www.archer.ac.uk/training/course-material/online/driving_test.php
- On successful completion, eligible users can apply for
 - account on ARCHER
 - 1,200 kAUs of time (80,000 core-hours) over 12 months
- Further information
 - Helpdesk: support@archer.ac.uk
 - Training: <http://www.archer.ac.uk/training/>.
- IPCC
 - <https://www.epcc.ed.ac.uk/research/computing/intel-parallel-computing-centre>

Women in HPC

- <http://www.womeninhpc.org.uk>



Working towards equal representation in HPC

Software sustainability institute

- <http://www.software.ac.uk/>



Software
Sustainability
Institute